
Computer Networks

Chap 6 Link Layer

P1



Computer Networking: A Top-Down Approach

8th edition, Global Edition

Jim Kurose, Keith Ross

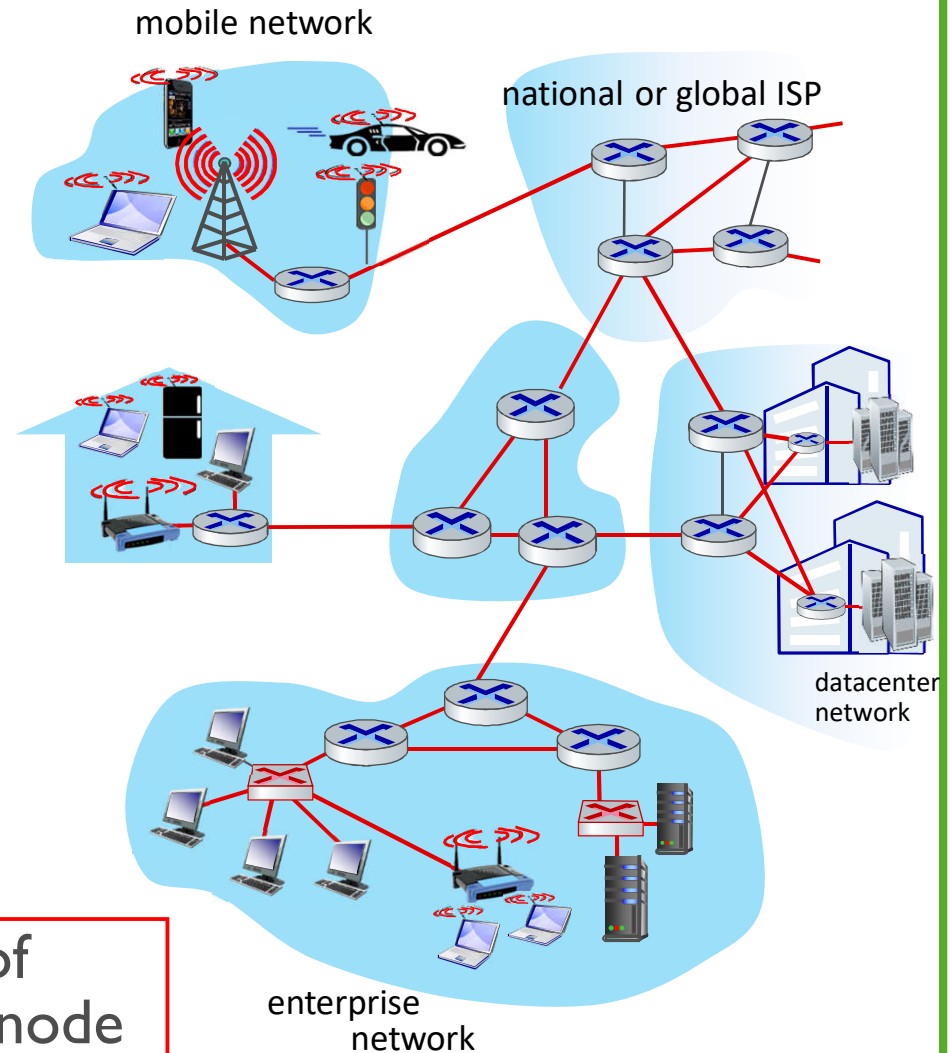
Copyright © 2022 Pearson Education Ltd

Link layer: introduction

terminology:

- ❖ **hosts and routers: nodes**
- ❖ **communication channels that connect adjacent nodes along communication path: links**
 - wired links
 - wireless links
 - LANs
- ❖ **layer-2 packet: frame, encapsulates datagram**

data-link layer has responsibility of transferring datagram from one node to *physically adjacent* node over a link



Link layer: context

❖ datagram transferred by different link protocols over different links:

- e.g., Ethernet on first link, frame relay on intermediate links, 802.11 on last link

❖ each link protocol provides different services

- e.g., may or may not provide rdt over link

transportation analogy:

❖ trip from Princeton to Lausanne

- limo: Princeton to JFK
- plane: JFK to Geneva
- train: Geneva to Lausanne

❖ **tourist = datagram**

❖ **transport segment = communication link**

❖ **transportation mode = link layer protocol**

❖ **travel agent = routing algorithm**

Link layer services

❖ *framing, link access:*

- encapsulate datagram into frame, adding header, trailer
- channel access if shared medium
- “MAC” addresses (Media Access Control) used in frame headers to identify source, destination
 - different from IP address!

❖ *reliable delivery between adjacent nodes*

- we learned how to do this already
- seldom used on low bit-error link (fiber, twisted pair)
- wireless links: high error rates
 - **Q:** why both link-level and end-end reliability?

Link layer services (more)

❖ **flow control:**

- pacing between adjacent sending and receiving nodes

❖ **error detection:**

- errors caused by signal attenuation, noise.
- receiver detects presence of errors:
 - signals sender for retransmission or drops frame

❖ **error correction:**

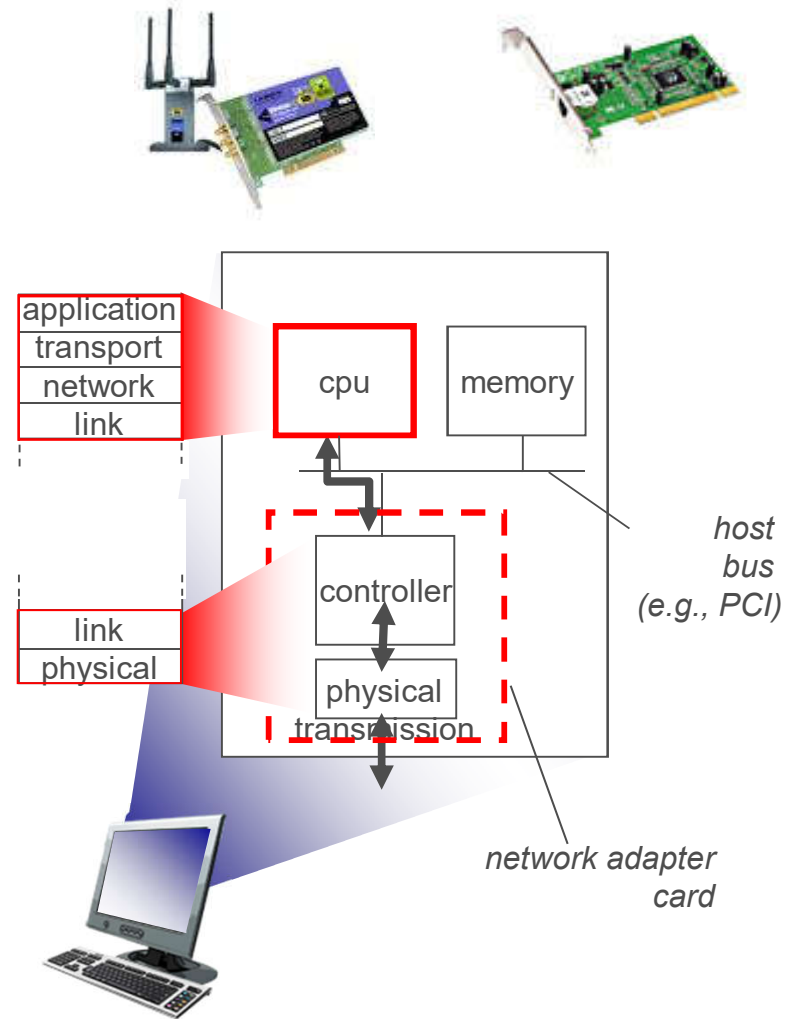
- receiver identifies *and corrects* bit error(s) without resorting to retransmission

❖ **half-duplex and full-duplex**


- with half duplex, nodes at both ends of link can transmit, but not at same time

Where is the link layer implemented?

- ❖ in each and every host
- ❖ link layer implemented in “adaptor” (aka **network interface card NIC**) or on a **chip**
 - Ethernet card, 802.11 card; Ethernet chipset
 - implements link, physical layer
- ❖ attaches into host's system buses
- ❖ combination of hardware, software, firmware



Adaptors communicating



❖ sending side:

- encapsulates datagram in frame
- adds error checking bits, reliable data transfer, flow control, etc.

❖ receiving side

- looks for errors, reliable data transfer, flow control, etc.
- extracts datagram, passes to upper layer at receiving side

Link layer, **LANs**: outline

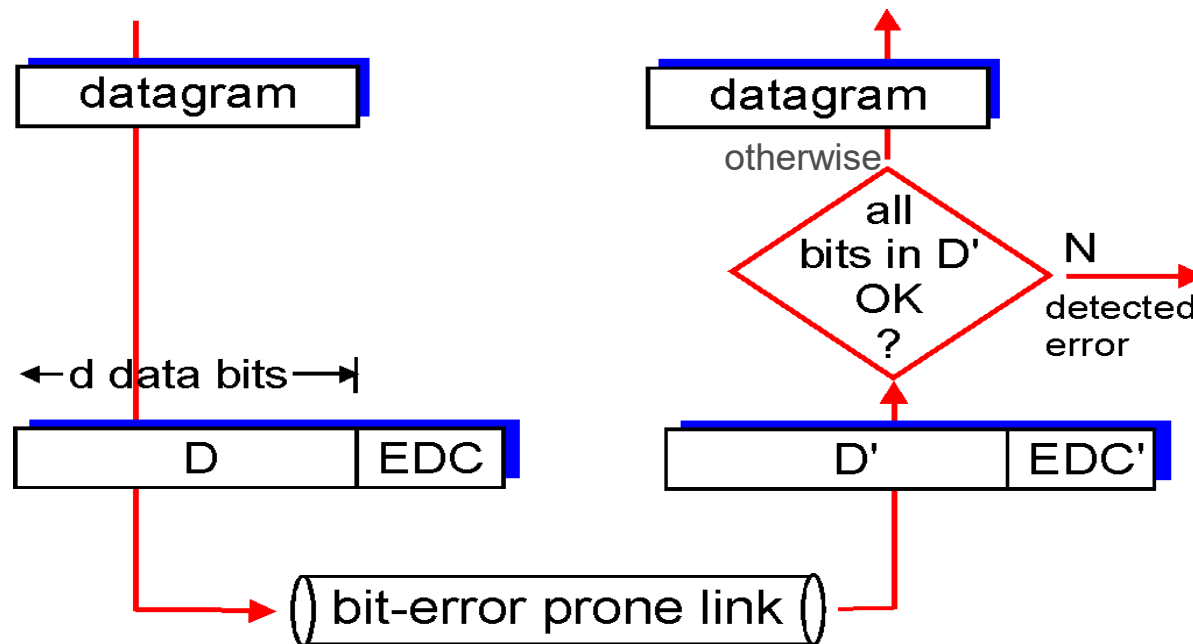
error detection, correction

Error detection

EDC= Error Detection and Correction bits (redundancy)

D = Data protected by error checking, may include header fields


- Error detection not 100% reliable!
 - protocol may miss some errors, but rarely
 - larger EDC field yields better detection and correction



Parity checking

single bit parity:

- detect single bit errors




Even/odd parity: set parity bit so there is an even/odd number of 1's

At receiver:

- compute parity of d received bits
- compare with received parity bit – if different than error detected

two-dimensional bit parity:

- detect and correct single bit errors



1	0	1	0	1	1
1	1	1	1	0	0
0	1	1	1	0	1
0	0	1	0	1	0

no errors

1	0	1	0	1	1
1	0	1	1	0	0
0	1	1	1	0	1
0	0	1	0	1	0

parity error

parity error

*correctable
single bit error*

Internet checksum (review)

goal: detect “errors” (e.g., flipped bits) in transmitted packet
(note: used at transport layer only)

sender:


- ❖ treat segment contents as sequence of 16-bit integers
- ❖ checksum: addition (1's complement sum) of segment contents
- ❖ sender puts checksum value into UDP checksum field

receiver:

- ❖ compute checksum of received segment
- ❖ check if computed checksum equals checksum field value:
 - NO - error detected
 - YES - no error detected.
But maybe errors nonetheless?

Cyclic redundancy check

- ❖ more powerful error-detection coding
- ❖ view data bits, **D**, as a binary number
- ❖ **G**: bit pattern (generator), of $r+1$ bits (given, specified in CRC standard)
- ❖ goal: choose r CRC bits, **R**, such that
 - $\langle D, R \rangle$ exactly divisible by G (modulo 2)
 - receiver knows G , divides $\langle D, R \rangle$ by G . If non-zero remainder: error detected!
 - can detect all burst errors less than $r+1$ bits



$$R = \text{remainder}\left[\frac{D \cdot 2^r}{G}\right]$$

Link layer, **LANs**: outline

multiple access protocols

Multiple access links, protocols

two types of “links”:

❖ point-to-point

- PPP for dial-up access
- point-to-point link between Ethernet switch, host

❖ **broadcast (shared wire or medium)**


- old-fashioned Ethernet
- upstream hybrid fiber-coaxial (HFC)
- 802.11 wireless LAN



shared wire (e.g.,
cabled Ethernet)



shared RF
(e.g., 802.11 WiFi)



shared RF
(satellite)



humans at a
cocktail party
(shared air, acoustical)

Multiple access protocols

- ❖ single shared broadcast channel
- ❖ two or more simultaneous transmissions by nodes: interference
 - *collision* if node receives two or more signals at the same time

multiple access protocol

- ❖ distributed algorithm that determines how nodes share channel, i.e., determine when node can transmit
- ❖ communication about channel sharing must use channel itself!
 - no out-of-band channel for coordination

An ideal multiple access protocol

given: broadcast channel of rate R bps

desiderata:

1. when one node wants to transmit, it can send at rate R .
2. when M nodes want to transmit, each can send at average rate R/M
3. fully decentralized:
 - no special node to coordinate transmissions
 - no synchronization of clocks, slots
4. simple

Media Access Control (MAC) protocols: taxonomy

three broad classes:

❖ *channel partitioning*

- divide channel into smaller “pieces” (time slots, frequency, code)
- allocate piece to node for exclusive use

❖ *random access*

- channel not divided, allow collisions
- “recover” from collisions


❖ *“taking turns”*

- nodes take turns, but nodes with more to send can take longer turns

Channel partitioning MAC protocols: TDMA

TDMA: time division multiple access


- ❖ access to channel in "rounds"
- ❖ each station gets fixed length slot (length = packet transmission time) in each round
- ❖ unused slots go idle
- ❖ example: 6-station LAN, 1,3,4 have packets to send, slots 2,5,6 idle



Channel partitioning MAC protocols: FDMA

FDMA: frequency division multiple access

- ❖ channel spectrum divided into frequency bands
- ❖ each station assigned fixed frequency band
- ❖ unused transmission time in frequency bands go idle
- ❖ example: 6-station LAN, 1,3,4 have packet to send, frequency bands 2,5,6 idle



Random access protocols

- ❖ **when node has packet to send**
 - transmit at full channel data rate R .
 - no *a priori* coordination among nodes
- ❖ **two or more transmitting nodes → “collision”,**
- ❖ **random access MAC protocol specifies:**
 - how to detect collisions
 - how to recover from collisions (e.g., via delayed retransmissions)
- ❖ **examples of random access MAC protocols:**
 - ALOHA, slotted ALOHA
 - CSMA, CSMA/CD, CSMA/CA


Slotted ALOHA

assumptions:


- ❖ all frames same size
- ❖ time divided into equal size slots (time to transmit 1 frame)
- ❖ nodes start to transmit only slot beginning
- ❖ nodes are synchronized
- ❖ if 2 or more nodes transmit in slot, all nodes detect collision

operation:

- ❖ when node obtains fresh frame, transmits in next slot
 - if no collision: node can send new frame in next slot
 - if collision: node retransmits frame in each subsequent slot with **prob.** p until success



Slotted ALOHA



Pros:


- ❖ single active node can continuously transmit at full rate of channel
- ❖ highly decentralized: only slots in nodes need to be in sync
- ❖ simple

Cons:

- ❖ collisions, wasting slots
- ❖ idle slots
- ❖ nodes may not be able to detect collision in less than time to transmit packet
- ❖ clock synchronization

Pure (unslotted) ALOHA

- ❖ unslotted Aloha: simpler, no synchronization
- ❖ when frame first arrives
 - transmit immediately
- ❖ collision probability increases:
 - frame sent at t_0 collides with other frames sent in $[t_0-1, t_0+1]$



CSMA (carrier sense multiple access)

CSMA: listen before
transmit:

if channel sensed idle:
transmit entire frame

❖ if channel sensed busy, defer
transmission

❖ human analogy: don't
interrupt others!

CSMA/CD (collision detection)

CSMA/CD: carrier sensing, deferral as in CSMA


- collisions *detected* within short time
- colliding transmissions aborted, reducing channel wastage

❖ **collision detection:**

- easy in wired LANs: measure signal strengths, compare transmitted, received signals
- difficult in wireless LANs: received signal strength overwhelmed by local transmission strength

❖ **human analogy: the polite conversationalist**

CSMA/CD (collision detection)



Ethernet CSMA/CD algorithm

1. NIC receives datagram from network layer, creates frame
2. If NIC senses channel idle, starts frame transmission. If NIC senses channel busy, waits until channel idle, then transmits.
3. If NIC transmits entire frame without detecting another transmission, NIC is done with frame !
4. If NIC detects another transmission while transmitting, aborts and sends jam signal
5. After aborting, NIC enters **binary (exponential) backoff**:
 - after m th collision, NIC chooses K at random from $\{0, 1, 2, \dots, 2^m - 1\}$. NIC waits $K \cdot 512$ bit times, returns to Step 2
 - longer backoff interval with more collisions

“Taking turns” MAC protocols

channel partitioning MAC protocols:

- share channel *efficiently* and *fairly* at high load
- inefficient at low load: delay in channel access, 1/N bandwidth allocated even if only 1 active node!

random access MAC protocols

- efficient at low load: single node can fully utilize channel
- high load: collision overhead


“taking turns” protocols

look for best of both worlds!

“Taking turns” MAC protocols

polling:

- ❖ master node
 - “invites” slave nodes to transmit in turn
- ❖ typically used with “dumb” slave devices
- ❖ concerns:
 - polling overhead
 - latency
 - single point of failure (master)



“Taking turns” MAC protocols

token passing:

- control *token* passed from one node to next sequentially.
- token message
- concerns:
 - token overhead
 - latency
 - single point of failure (token)

